

```

%pip install llama_index ftfy regex tqdm
%pip install llama-index-vector-stores-qdrant llama-index-embeddings-clip
%pip install git+https://github.com/openai/CLIP.git
%pip install torch torchvision
%pip install matplotlib scikit-image
%pip install qdrant_client
%pip install pymupdf
%pip install qdrant_client
%pip install llama_index.embeddings.huggingface

import os
import fitz # PyMuPDF
from PIL import Image
import io
import hashlib

import qdrant_client
from llama_index.vector_stores.qdrant import QdrantVectorStore
from llama_index.core import StorageContext, Settings, SimpleDirectoryReader
from llama_index.core.indices import MultiModalVectorStoreIndex
from llama_index.embeddings.huggingface import HuggingFaceEmbedding

# -----
# Utility: deterministic file name
# -----
def make_filename(base_name, page, suffix, ext="jpg"):
    return f"{os.path.splitext(base_name)[0]}_p{page}_{suffix}.{ext}"

# -----
# Extraction Pipeline
# -----
def extract_images(pdf_path, method="both", dpi=200):
    """
    Extract images from PDF.
    method = "embedded" | "fullpage" | "both"
    Saves images into /artifacts/embedded and /artifacts/fullpage
    """
    pdf_document = fitz.open(pdf_path)

    file_name = os.path.basename(pdf_path)
    pdf_dir = os.path.dirname(pdf_path)

    # Output dirs
    artifact_dir = os.path.join(pdf_dir, "artifacts")
    embedded_dir = os.path.join(artifact_dir, "embedded")
    fullpage_dir = os.path.join(artifact_dir, "fullpage")

    if method in ["embedded", "both"]:
        os.makedirs(embedded_dir, exist_ok=True)
    if method in ["fullpage", "both"]:
        os.makedirs(fullpage_dir, exist_ok=True)

```

```

for page_num, page in enumerate(pdf_document):
    page_number = page_num + 1

    # --- Embedded images ---
    if method in ["embedded", "both"]:
        for img_index, img in enumerate(page.get_images(full=True)):
            xref = img[0]
            base_image = pdf_document.extract_image(xref)
            image_bytes = base_image["image"]
            image_ext = base_image["ext"]

            filename = make_filename(file_name, page_number, f"emb{img_index+
            filepath = os.path.join(embedded_dir, filename)
            with open(filepath, "wb") as f:
                f.write(image_bytes)

    # --- Full page render ---
    if method in ["fullpage", "both"]:
        zoom = dpi / 72
        mat = fitz.Matrix(zoom, zoom)
        pix = page.get_pixmap(matrix=mat)
        img_bytes = pix.tobytes("png")
        img = Image.open(io.BytesIO(img_bytes))

        filename = make_filename(file_name, page_number, "page", "jpg")
        page_img_path = os.path.join(fullpage_dir, filename)
        img.convert("RGB").save(page_img_path, "JPEG", quality=90)

pdf_document.close()

# -----
# Qdrant Setup
# -----
embed_model_text = HuggingFaceEmbedding(model_name="BAAI/bge-small-en-v1.5")
embed_model_image = HuggingFaceEmbedding(model_name="sentence-transformers/clip-V

Settings.chunk_size = 512 # no global embed_model
Settings.embed_model = embed_model_image

client = qdrant_client.QdrantClient(path="qdrant_db_test")

text_store = QdrantVectorStore(
    client=client,
    collection_name="text_collection",
    embed_model=embed_model_text,
)

image_store = QdrantVectorStore(
    client=client,
    collection_name="image_collection",
    embed_model=embed_model_image,
)

```

```

storage_context = StorageContext.from_defaults(
    vector_store=text_store,
    image_store=image_store,
)

# -----
# Main function: add PDF to index
# -----
def add_pdf_to_index(doc_dir, pdf_path, method="both"):
    # Step 1: Extract images
    extract_images(pdf_path, method=method)

    # Step 2: Use SimpleDirectoryReader to load extracted artifacts
    pdf_dir = os.path.dirname(pdf_path)
    embedded_artifact_dir = os.path.join(pdf_dir, "artifacts/embedded")
    fullpage_artifact_dir = os.path.join(pdf_dir, "artifacts/fullpage")

    documents = SimpleDirectoryReader(doc_dir).load_data()
    if os.path.exists(embedded_artifact_dir):
        documents += SimpleDirectoryReader(embedded_artifact_dir).load_data()
    if os.path.exists(fullpage_artifact_dir):
        documents += SimpleDirectoryReader(fullpage_artifact_dir).load_data()

    if not documents:
        print(f"No documents found for {pdf_path}")
        return None

    # Step 3: Insert into multimodal index
    index = MultiModalVectorStoreIndex.from_documents(
        documents,
        storage_context=storage_context,
        show_progress=True,
    )
    return index

# Upload a file from frontend, then make a folder for that pdf and put it in ther
doc_dir = "data_fin/Duravent"
pdf_path = "data_fin/Duravent/DURAVENT_SK_LD_CT-001E_02-24_v3.pdf"
ll_index = add_pdf_to_index(doc_dir, pdf_path, method="fullpage")

doc_dir = "data_fin/Nvidia"
pdf_path2 = "data_fin/Nvidia/NVIDIA-2025-Annual-Report.pdf"
ll_index = add_pdf_to_index(doc_dir, pdf_path2, method="embedded")

print(client.get_collections())

from llama_index.core.query_engine import CustomQueryEngine
from llama_index.core.retrievers import BaseRetriever
from llama_index.core.schema import ImageNode, NodeWithScore, MetadataMode, T
from llama_index.core.prompts import PromptTemplate
from llama_index.core.base.response.schema import Response

```

```

from google import genai
from google.genai import types
from typing import Optional, List
import base64
from pydantic import BaseModel, Field
import json
from google.colab import userdata

```

```
llm_client = genai.Client(api_key=userdata.get("GOOGLE_API_KEY"))
```

```
QA_PROMPT_TMPL = """Below we give parsed text and images as context.
```

```
Use both the parsed text and images to answer the question.
```

```
Write your response in markdown
```

```
-----
{context_str}
-----
```

```
Given the context information and not prior knowledge, answer the query. Explain
from the text or image, and if there's discrepancies, and your reasoning for
```

```
Query: {query_str}
```

```
Answer: """
```

```
QA_PROMPT = PromptTemplate(QA_PROMPT_TMPL)
```

```
class Structure(BaseModel):
```

```
    text_response: str = Field(description="Text response from the LLM")
```

```
    file_name_used: List[str] = Field(description="List of image nodes file_nam
```

```
class MultimodalGeminiEngine(CustomQueryEngine):
```

```
    """More robust version with better error handling and image format detect:
```

```
    qa_prompt: PromptTemplate
```

```
    retriever: BaseRetriever
```

```
    def __init__(self, qa_prompt: Optional[PromptTemplate] = None, **kwargs) -
```

```
        """Initialize."""
```

```
        super().__init__(qa_prompt=qa_prompt or QA_PROMPT, **kwargs)
```

```
    def _get_image_mime_type(self, image_path: str) -> str:
```

```
        """Determine MIME type from file extension."""
```

```
        extension = image_path.lower().split('.')[-1]
```

```
        mime_types = {
```

```
            'jpg': 'image/jpeg',
```

```
            'jpeg': 'image/jpeg',
```

```
            'png': 'image/png',
```

```
            'gif': 'image/gif',
```

```
            'webp': 'image/webp',
```

```
            'bmp': 'image/bmp'
```

```
        }
```

```
        return mime_types.get(extension, 'image/jpeg') # Default to JPEG
```

```

def _process_image_node(self, image_node: ImageNode) -> Optional[types.Part]:
    """Process a single image node into a GenAI Part."""

    # Try each method in order of preference
    methods = [
        self._try_base64_image,
        self._try_resolve_image,
        self._try_file_path,
        self._try_image_url
    ]

    for method in methods:
        try:
            part = method(image_node)
            if part is not None:
                return part
        except Exception as e:
            continue

    print(f"Warning: Could not process ImageNode {image_node.id_}")
    return None

def _try_base64_image(self, image_node: ImageNode) -> Optional[types.Part]:
    """Try to get image from base64 encoded data."""
    try:
        if hasattr(image_node, 'image') and image_node.image:
            image_bytes = base64.b64decode(image_node.image)
            return types.Part.from_bytes(data=image_bytes, mime_type="image/jpeg")
        return None
    except Exception as e:
        return f"ERROR : {e}"

def _try_resolve_image(self, image_node: ImageNode) -> Optional[types.Part]:
    """Try to get image using resolve_image method."""
    if hasattr(image_node, 'resolve_image'):
        image_buffer = image_node.resolve_image()
        image_bytes = image_buffer.getvalue()
        return types.Part.from_bytes(data=image_bytes, mime_type="image/jpeg")
    return None

def _try_file_path(self, image_node: ImageNode) -> Optional[types.Part]:
    """Try to get image from file path."""
    if hasattr(image_node, 'image_path') and image_node.image_path:
        with open(image_node.image_path, 'rb') as f:
            image_bytes = f.read()
            mime_type = self._get_image_mime_type(image_node.image_path)
            return types.Part.from_bytes(data=image_bytes, mime_type=mime_type)
    return None

def _try_image_url(self, image_node: ImageNode) -> Optional[types.Part]:
    """Try to get image from URL."""
    if hasattr(image_node, 'image_url') and image_node.image_url:
        # For cloud storage URLs (gs://) or public URLs
        return types.Part.from_uri(
            file_uri=image_node.image_url,

```

```

        mime_type="image/jpeg"
    )
    return None

def custom_query(self, query_str: str):
    """Execute the query with robust image processing."""
    # Retrieve nodes
    nodes = self.retriever.retrieve(query_str)
    img_nodes = [n for n in nodes if isinstance(n.node, ImageNode)]
    text_nodes = [n for n in nodes if isinstance(n.node, TextNode)]

    # Create context string
    context_str = "\n\n".join(
        [r.get_content(metadata_mode=MetadataMode.LLM) for r in nodes]
    )

    fmt_prompt = self.qa_prompt.format(context_str=context_str, query_str:

    # Prepare content parts
    content_parts = [fmt_prompt]

    # Process image nodes
    successful_images = 0
    for img_node in img_nodes:
        image_part = self._process_image_node(img_node.node)
        if image_part:
            content_parts.append(img_node.get_content(metadata_mode=Metadi
            content_parts.append(image_part)
            successful_images += 1

    print(f"Successfully processed {successful_images}/{len(img_nodes)} in

    try:
        # Generate content with structured output
        response = llm_client.models.generate_content(
            model="gemini-2.5-flash",
            contents=content_parts,
            config={
                "response_mime_type": "application/json",
                "response_schema": Structure,
            }
        )

        structured_response = response.parsed
        full_response = f"{structured_response.text_response}\n\nImages U:

    except Exception as e:
        print(f"Structured output failed: {e}")
        # Fallback to regular response
        try:
            response = client.models.generate_content(
                model="gemini-2.5-flash",
                contents=content_parts,
            )
            full_response = response.text
            structured_response = None

```

```
        structured_response = None
    except Exception as fallback_error:
        print(f"Fallback failed: {fallback_error}")
        full_response = f"Error: {str(e)}"
        structured_response = None

    return Response(
        response=full_response,
        source_nodes=nodes,
        metadata={
            "text_nodes": text_nodes,
            "image_nodes": img_nodes,
            "successful_images": successful_images,
            "total_images": len(img_nodes),
            "structured_response": structured_response
        }
    )

retriever = ll_index.as_retriever(similarity_top_k=70, image_similarity_top_k=10)

engine = MultimodalGeminiEngine(retriever=retriever)

response = engine.custom_query("What was nvidia's perfomance in this year")

from IPython.display import Markdown, display
resp = str(response)
display(Markdown(resp))
```